

broom: An R Package to Convert Statistical Models into Tidy Data Frames

David Robinson
4/11/2015

What is tidy data?


Data frames arranged as:

- One row for each *observation*
- One column for each *variable*
- One table for each *type of observational unit*

For details, see [Tidy Data \(Wickham 2014\)](#)


“Tidy tools” work with tidy data frames

tidyverse




reshape data to be tidy

dplyr



manipulate and summarize tidy data


ggplot2




visualize tidy data

Source: RStudio Data Wrangling Cheatsheet
RStudio Data Visualization Cheatsheet


Tidy tools work together in exploratory data analysis



Everything works well until...



Visualizing and manipulating model objects is difficult



**Model objects are
messy**

Example:
linear regression

What's “messy” about a linear regression?

```
> lmfit <- lm(mpg ~ wt + qsec, mtcars)
```

What's “messy” about a linear regression?

```
> summary(lmfit)
```

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

Min	10	Median	30	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	19.7462	5.2521	3.760	0.000765 ***
wt	-5.0480	0.4840	-10.430	2.52e-11 ***
qsec	0.9292	0.2650	3.506	0.001500 **

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 2.596 on 29 degrees of freedom

Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144

F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

What's “messy” about a linear regression?

```
> summary(lmfit)
```

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

Min	10	Median	30	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

1. Extracting coefficients takes multiple steps:

```
data.frame(coef(summary(lmfit)))
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	19.7462	5.2521	3.760	0.000765	***
wt	-5.0480	0.4840	-10.430	2.52e-11	***
qsec	0.9292	0.2650	3.506	0.001500	**

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 2.596 on 29 degrees of freedom

Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144

F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

What's “messy” about a linear regression?

```
> summary(lmfit)
```

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

Min	10	Median	30	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	19.7462	5.2521	3.760	0.000765	***
wt	-5.0480	0.4840	-10.430	2.52e-11	***
qsec	0.9292	0.2650	3.506	0.001500	**

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 2.596 on 29 degrees of freedom

Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144

F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

2. Information stored in row names
(can't combine models)

What's “messy” about a linear regression?

```
> summary(lmfit)
```

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

Min	10	Median	30	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

3. Column names are inconvenient (access with \$“Pr(>|t|)”, converts to Pr...t...)

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	19.7462	5.2521	3.760	0.000765	***
wt	-5.0480	0.4840	-10.430	2.52e-11	***
qsec	0.9292	0.2650	3.506	0.001500	**

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 2.596 on 29 degrees of freedom

Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144

F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

What's “messy” about a linear regression?

```
> summary(lmfit)
```

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

Min	10	Median	30	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	19.7462	5.2521	3.760	0.000765 ***
wt	-5.0480	0.4840	-10.430	2.52e-11 ***
qsec	0.9292	0.2650	3.506	0.001500 **

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 2.596 on 29 degrees of freedom

Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144

F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

What's “messy” about a linear regression?

```
> summary(lmfit)
```

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

Min	10	Median	30	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

1. Extracting coefficients takes multiple steps:

```
data.frame(coef(summary(lmfit)))
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	19.7462	5.2521	3.760	0.000765	***
wt	-5.0480	0.4840	-10.430	2.52e-11	***
qsec	0.9292	0.2650	3.506	0.001500	**

2. Information stored in row names (can't combine models)

3. Column names are inconvenient and inconsistent

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 2.596 on 29 degrees of freedom

Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144

F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

4. Information is computed in print method, not stored

**These inconveniences
aren't an exception,
they're the rule**

broom's `tidy()` method does the work for you

```
> tidy(lmfit)
  term estimate std.error statistic p.value
1 (Intercept) 19.746     5.252      3.76 7.65e-04
2          wt -5.048     0.484     -10.43 2.52e-11
3         qsec  0.929     0.265      3.51 1.50e-03
```

broom's `tidy()` method does the work for you

One function
to call


```
> tidy(lmfit)
```

	term	estimate	std.error	statistic	p.value
1	(Intercept)	19.746	5.252	3.76	7.65e-04
2	wt	-5.048	0.484	-10.43	2.52e-11
3	qsec	0.929	0.265	3.51	1.50e-03

Convenient
column names

Information stored
in columns, never
row names

broom takes model objects and
turns them into tidy data frames
that can be used with tidy tools



Introduction to broom

```
> install.packages("broom")
> library(broom)
```

broom's three methods

- broom defines tidying methods for extracting three kinds of statistics from an object:
 - **tidy()**: component-level statistics
 - **augment()**: observation-level statistics
 - **glance()**: model-level statistics

Example: three levels of a linear regression

```
> summary(lmfit)
```

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

Min	10	Median	30	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

Observation Level:
fitted values, residuals
augment()

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	19.7462	5.2521	3.760	0.000765	***
wt	-5.0480	0.4840	-10.430	2.52e-11	***
qsec	0.9292	0.2650	3.506	0.001500	**

Component Level:
coefficients, p-values
tidy()

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Model Level:
 R^2 , F-statistic,
deviance
glance()

Residual standard error: 2.596 on 29 degrees of freedom
Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144
F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

The `tidy()` method: component-level statistics

```
> tidy(lmfit)
  term estimate std.error statistic p.value
1 (Intercept) 19.746     5.252      3.76 7.65e-04 ← each row is a coefficient
2          wt -5.048     0.484     -10.43 2.52e-11
3         qsec  0.929     0.265      3.51 1.50e-03
```

The augment() method: observation-level statistics

```
> augment(lmfit)
```

	.rownames	mpg	wt	qsec	.fitted	.se.fit	.resid	.hat	.sigma
1	Mazda RX4	21.0	2.62	16.5	21.82	0.683	-0.8151	0.0693	2.64
2	Mazda RX4 Wag	21.0	2.88	17.0	21.05	0.547	-0.0482	0.0444	2.64
3	Datsun 710	22.8	2.32	18.6	25.33	0.640	-2.5273	0.0607	2.60
4	Hornet 4 Drive	21.4	3.21	19.4	21.58	0.623	-0.1806	0.0576	2.64
5	Hornet Sportabout	18.7	3.44	17.0	18.20	0.512	0.5039	0.0389	2.64
6	Valiant	18.1	3.46	20.2	21.07	0.803	-2.9686	0.0957	2.58
7	Duster 360	14.3	3.57	15.8	16.44	0.701	-2.1434	0.0729	2.61
8	Merc 240D	24.4	3.19	20.0	22.23	0.730	2.1729	0.0791	2.61
9	Merc 230	22.8	3.15	22.9	25.12	1.410	-2.3237	0.2950	2.59
10	Merc 280	19.2	3.44	18.3	19.39	0.491	-0.1855	0.0358	2.64
11	Merc 280C	17.8	3.44	18.9	19.94	0.557	-2.1430	0.0460	2.61
12	Merc 450SE	16.4	4.07	17.4	15.37	0.615	1.0310	0.0561	2.63
13	Merc 450SL	17.3	3.73	17.6	17.27	0.520	0.0289	0.0402	2.64
14	Merc 450SLC	15.2	3.78	18.0	17.39	0.539	-2.1904	0.0431	2.61
15	Cadillac Fleetwood	10.4	5.25	18.0	9.95	1.092	0.4487	0.1768	2.64
16	Lincoln Continental	10.4	5.42	17.8	8.92	1.161	1.4757	0.2001	2.62
17	Chrysler Imperial	14.7	5.34	17.4	8.95	1.115	5.7486	0.1844	2.35
18	Fiat 128	32.4	2.20	19.5	26.73	0.751	5.6679	0.0836	2.39
19	Honda Civic	30.4	1.61	18.5	28.80	0.892	1.5975	0.1180	2.62
20	Toyota Corolla	33.9	1.83	19.9	28.97	0.909	4.9258	0.1226	2.45

each row is an observation from the original data

The augment() method: observation-level statistics

```
> augment(lmfit)          note that new columns start with .  
  .rownames  mpg   wt  qsec .fitted .se.fit .resid   .hat .sigma  
1  Mazda RX4  21.0 2.62 16.5  21.82  0.683 -0.8151 0.0693 2.64  
2  Mazda RX4 Wag 21.0 2.88 17.0  21.05  0.547 -0.0482 0.0444 2.64  
3  Datsun 710 22.8 2.32 18.6  25.33  0.640 -2.5273 0.0607 2.60  
4  Hornet 4 Drive 21.4 3.21 19.4  21.58  0.623 -0.1806 0.0576 2.64  
5  Hornet Sportabout 18.7 3.44 17.0  18.20  0.512  0.5039 0.0389 2.64  
6  Valiant 18.1 3.46 20.2  21.07  0.803 -2.9686 0.0957 2.58  
7  Duster 360 14.3 3.57 15.8  16.44  0.701 -2.1434 0.0729 2.61  
8  Merc 240D 24.4 3.19 20.0  22.23  0.730  2.1729 0.0791 2.61  
9  Merc 230 22.8 3.15 22.9  25.12  1.410 -2.3237 0.2950 2.59  
10  Merc 280 19.2 3.44 18.3  19.39  0.491 -0.1855 0.0358 2.64  
11  Merc 280C 17.8 3.44 18.9  19.94  0.557 -2.1430 0.0460 2.61  
12  Merc 450SE 16.4 4.07 17.4  15.37  0.615  1.0310 0.0561 2.63  
13  Merc 450SL 17.3 3.73 17.6  17.27  0.520  0.0289 0.0402 2.64  
14  Merc 450SLC 15.2 3.78 18.0  17.39  0.539 -2.1904 0.0431 2.61  
15 Cadillac Fleetwood 10.4 5.25 18.0  9.95  1.092  0.4487 0.1768 2.64  
16 Lincoln Continental 10.4 5.42 17.8  8.92  1.161  1.4757 0.2001 2.62  
17 Chrysler Imperial 14.7 5.34 17.4  8.95  1.115  5.7486 0.1844 2.35  
18 Fiat 128 32.4 2.20 19.5  26.73  0.751  5.6679 0.0836 2.39  
19 Honda Civic 30.4 1.61 18.5  28.80  0.892  1.5975 0.1180 2.62  
20 Toyota Corolla 33.9 1.83 19.9  28.97  0.909  4.9258 0.1226 2.45
```

each row is an observation from the original data

The `glance()` method: model-level statistics

```
> glance(lmfit)
   r.squared adj.r.squared sigma statistic p.value df logLik AIC BIC deviance
1    0.826          0.814   2.6       69 9.39e-12   3 -74.4 157 163      195 ← one row
                                         for the model
```

broom works across
many kinds of model
objects

Nonlinear least squares: before

```
> n <- nls(mpg ~ k * e ^ wt, data = mtcars, start = list(k = 1, e = 2))  
> summary(n)
```

Formula: mpg ~ k * e^wt

Parameters:

	Estimate	Std. Error	t value	Pr(> t)
k	49.6597	3.7888	13.1	6e-14 ***
e	0.7456	0.0199	37.5	<2e-16 ***

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 2.67 on 30 degrees of freedom

Number of iterations to convergence: 10

Achieved convergence tolerance: 2.04e-06

Nonlinear least squares: after

K-means clustering: **before**

```
> k  
K-means clustering with 3 clusters of sizes 47, 103, 100
```

Cluster means:

	oracle	x1	x2
1	3.00	-3.3256	-2.398
2	2.03	0.0477	0.901
3	1.00	4.8963	-1.188

Clustering vector:

Within cluster sum of squares by cluster:

[1] 81.5 206.4 181.4
(between SS / total SS = 86.5 %)

K-means clustering: after



```
> tidy(k)
      x1      x2      x3 size withinss cluster
1 3.00 -3.3256 -2.398   47     81.5       1 ← each row is one
2 2.03  0.0477  0.901  103    206.4       2
3 1.00  4.8963 -1.188  100    181.4       3
> augment(k, kdat)
  oracle      x1      x2 .cluster
1      2  0.345  1.512       2 ← each row is one
2      1  5.784  0.246       3
3      2 -0.291  1.378       2
4      2 -0.922  0.503       2
5      2 -0.456  0.860       2
6      2 -0.897  1.247       2
...
> glance(k)
  totss tot.withinss betweenss iter
1 3484        469        3015    2 ← one row describing the
                                entire clustering
                                operation
```

And many others...

package	class	tidy	augment	glance
base	<code>data.frame</code> , <code>matrix</code>	✓		✓
	<code>table</code>	✓		
stats	<code>anova</code> , <code>aov</code> , <code>aovlist</code> , <code>density</code> , <code>ftable</code> , <code>manova</code> , <code>pairwise.htest</code> , <code>power.htest</code> , <code>spec</code> , <code>ts</code> , <code>TukeyHSD</code>	✓		
	<code>kmeans</code> , <code>lm</code> , <code>glm</code> , <code>nls</code>	✓	✓	✓
	<code>smooth.spline</code>		✓	✓
	<code>Arima</code> , <code>htest</code> , <code>summaryDefault</code>	✓		✓
biglm	<code>biglm</code> , <code>bigglm</code>	✓		✓
glmnet	<code>cv.glmnet</code> , <code>glmnet</code>	✓		✓
lfe	<code>felm</code>	✓	✓	✓
lmtest	<code>coeftest</code>	✓	✓	✓
lme4	<code>mer</code> , <code>merMod</code>	✓	✓	✓
maps	<code>map</code>	✓		
MASS	<code>ridgelm</code>	✓		✓
multcomp	<code>cld</code> , <code>confint.glht</code> , <code>glht</code> , <code>summary.glht</code>	✓		
plm	<code>plm</code>	✓	✓	✓
sp	<code>Line</code> , <code>Lines</code> , <code>Polygon</code> , <code>Polygons</code> , <code>SpatialLinesDataFrame</code> , <code>SpatialPolygons</code> , <code>SpatialPolygonsDataFrame</code>	✓		
survival	<code>aareg</code> , <code>cch</code> , <code>pyears</code> , <code>survexp</code> , <code>survfit</code>	✓		✓
	<code>coxph</code> , <code>survreg</code>	✓	✓	✓
zoo	<code>zoo</code>	✓		

Why are tidy models
useful?


ggplot2 can visualize tidy data



Source: RStudio Data Visualization Cheatsheet


Example: coefficient plot

```
td <- tidy(lmfit, conf.int = TRUE)
ggplot(td, aes(estimate, term, color = term)) +
  geom_point() +
  geom_errorbarh(aes(xmin = conf.low, xmax = conf.high))
```



Example: survival curves


```
library(survival)
surv_fit <- survfit(coxph(Surv(time, status) ~ age + sex, lung))
td <- tidy(surv_fit)
ggplot(td, aes(time, estimate)) + geom_line() +
  geom_ribbon(aes(ymin = conf.low, ymax = conf.high), alpha = .2)
```




Example: LASSO regression

```
tidied_cv <- tidy(glmnet_fit)
glance_cv <- glance(glmnet_fit)

ggplot(tidied_cv, aes(lambda, estimate)) + geom_line(color = "red") +
  geom_ribbon(aes(ymin = conf.low, ymax = conf.high), alpha = .2) +
  scale_x_log10() +
  geom_vline(xintercept = glance_cv$lambda.min) +
  geom_vline(xintercept = glance_cv$lambda.1se, lty = 2)
```




Tidy models can be combined and compared




	term	estimate	std.error	statistic	p.value
1	(Intercept)	19.746	5.252	3.76	7.65e-04
2	wt	-5.048	0.484	-10.43	2.52e-11
3	qsec	0.929	0.265	3.51	1.50e-03

- different parameters
- different methods
- bootstrap replicates
- subgroup models (within each country, gene...)
- ensemble voting

Tidy models can be combined and compared



Tidy models can be combined and compared




	term	estimate	std.error	statistic	p.value
1	(Intercept)	19.746	5.252	3.76	7.65e-04
2	wt	-5.048	0.484	-10.43	2.52e-11
3	qsec	0.929	0.265	3.51	1.50e-03

- different parameters
- different methods
- bootstrap replicates
- subgroup models (within each country, gene...)
- ensemble voting


If you can plot one nonlinear least squares fit...

```
augmented <- augment(nlsfit)  
ggplot(augmented, aes(wt, .fitted)) + geom_line()
```




...you can plot 50 bootstrap replicates of it

```
ggplot(combined_augmented, aes(wt, .fitted, group = replicate)) +  
  geom_line(alpha = .2)
```




If you can plot one instance of k-means clustering...

```
ggplot(assignments, aes(x1, x2)) +  
  geom_point(aes(color = .cluster)) +  
  geom_point(data = clusters, size = 10, shape = "X")
```



...you can plot it for many values of **k**

```
ggplot(combined_assignments, aes(x1, x2)) +  
  geom_point(aes(color = .cluster)) +  
  geom_point(data = combined_clusters, size = 10, shape = "X") +  
  facet_wrap(~ k)
```



Learn more: vignettes

[Introduction to broom](#)

[broom and dplyr](#)

[kmeans with dplyr+broom](#)

[Tidy bootstrapping with dplyr+broom](#)

Learn more: manuscript

<http://arxiv.org/pdf/1412.3565v2.pdf>

broom: An R Package for Converting Statistical Analysis Objects Into Tidy Data Frames

David Robinson

Abstract

The concept of "tidy data" offers a powerful framework for structuring data to ease manipulation, modeling and visualization. However, most R functions, both those built-in and those found in third-party packages, produce output that is not tidy, and that is therefore difficult to reshape, recombine, and otherwise manipulate. Here I introduce the **broom** package, which turns the output of model objects into tidy data frames that are suited to further analysis, manipulation, and visualization with input-tidy tools. **broom** defines the **tidy**, **augment**, and **glance** generics, which arrange a model into three levels of tidy output respectively: the component level, the observation level, and the model level. I provide examples to demonstrate how these generics work with tidy tools to allow analysis and modeling of data that is divided into subsets, to recombine results from bootstrap replicates, and to perform simulations that investigate the effect of varying input parameters.

Learn more: GitHub

<https://github.com/dgrtwo/broom>

 [dgrtwo / broom](#)

[Watch 17](#) [Star 199](#) [Fork 28](#)

Convert statistical analysis objects from R into tidy format

 102 commits  1 branch  7 releases  4 contributors

  branch: **master** [+ broom](#)

Added tidier for return value of optim, which is called from a tidy.l... 

 dgrtwo	authored 3 days ago	latest commit  9995e045ef
 R	Added tidier for return value of optim, which is called from a tidy.l...	3 days ago
 man-roxygen	Overhaul of how augmenting works across many objects. In particular t...	5 months ago
 man	Added tidier for return value of optim, which is called from a tidy.l...	3 days ago
 tests	Removed warning for lm fits when link function is log; added tests th...	4 months ago
 vignettes	Added `gam` to README. Removed rownames from glmnet output. Few typo ...	4 months ago
 .Rbuildignore	Update cran comments.	3 months ago
 .gitignore	Update cran comments.	3 months ago
 DESCRIPTION	Added `tidy` and `glance` for "biglm" and "bigglm" objects from the b...	23 days ago

Code

 Issues 9

 Pull requests 1

Pulse

Graphs

HTTPS clone URL

<https://github.com/dgrtwo/broom> 

You can clone with [HTTPS](#) or [Subversion](#). 

 [Clone in Desktop](#)

 [Download ZIP](#)

Contribute!

<https://github.com/dgrtwo/broom/issues>

 [dgrtwo / broom](#)

[Watch 17](#) [Star 199](#) [Fork 28](#)

Convert statistical analysis objects from R into tidy format

102 commits 1 branch 7 releases 4 contributors

branch: master → [broom / +](#)

Added tidier for return value of optim, which is called from a tidy.l... [...](#)

 dgrtwo authored 3 days ago latest commit 9995e045ef [Edit](#)

	Added tidier for return value of optim, which is called from a tidy.l...	3 days ago
	Overhaul of how augmenting works across many objects. In particular t...	5 months ago
	Added tidier for return value of optim, which is called from a tidy.l...	3 days ago
	Removed warning for lm fits when link function is log; added tests th...	4 months ago
	Added `gam` to README. Removed rownames from glmnet output. Few typo ...	4 months ago
	Update cran comments.	3 months ago
	Update cran comments.	3 months ago
	Added `tidy` and `glance` for "biglm" and "bigglm" objects from the b...	23 days ago

Code

- [Issues 9](#)
- [Pull requests 1](#)

[Pulse](#)

[Graphs](#)

HTTPS clone URL

<https://github.com/dgrtwo/broom.git> [Edit](#)

You can clone with [HTTPS](#) or [Subversion](#). [?](#)

[Clone in Desktop](#)

[Download ZIP](#)

Thank you!

- broom package
 - Matthieu Gomez
 - Boris Demeshev
 - Hadley Wickham
- Presentation
 - Dima Gorenstein
 - Storey Lab at Princeton University
 - UP-STAT 2015